

script_glm_thr.R

sergio

2020-05-27

```
library(glmnet)

## Loading required package: Matrix
## Loaded glmnet 3.0-2
source('../utils/utils_oblig.R')

## Loading required package: caret
## Loading required package: lattice
## Loading required package: ggplot2
## Registered S3 methods overwritten by 'ggplot2':
##   method      from
##   [.quosures    rlang
##   c.quosures    rlang
##   print.quosures rlang
set.seed(117)

script.name <- 'glm_thr'

script.date <- date()

script.start <- Sys.time()

print('Start')

## [1] "Start"
# leer el archivo dataset.csv de la carpeta

dataset <- read.csv('../data/dataset.csv')

# ver la estructura del dataset

# str(dataset)

# asignar el nombre del jugador como nombre de la fila

rownames(dataset) <- dataset$CustomerID

df <- na.omit(dataset[,-1])

df$ServiceArea <- NULL

print('** Distribucion a-priori de la variable a predecir')

## [1] "** Distribucion a-priori de la variable a predecir"
```

```

print(prop.table(table(df$Churn)))

##
##          No        Yes
## 0.7131871 0.2868129

df.part <- train_dev_partition(df, p = 0.9)

df.thr_vec <- seq(0.1, 0.9, 0.05)

df.fn_summary <- function(data, lev = NULL, model = NULL) {
  fn_summaryUtilityThr(data, df.thr_vec)
}

df.metric <- 'utility'

df.form <- Churn ~ .

print('** GLM')

## [1] "** GLM"

df.glm.ctrl <- trainControl(method = 'cv',
                             number = 5,
                             verboseIter = TRUE,
                             classProbs = TRUE,
                             search = 'random',
                             summaryFunction = df.fn_summary)

df.glm <- train(form = df.form,
                 data = df.part$train,
                 method = 'glmnet',
                 family = 'binomial',
                 trControl = df.glm.ctrl,
                 tuneLength = 5,
                 metric = df.metric)

## + Fold1: alpha=0.26189, lambda=0.10504
## - Fold1: alpha=0.26189, lambda=0.10504
## + Fold1: alpha=0.17301, lambda=3.96064
## - Fold1: alpha=0.17301, lambda=3.96064
## + Fold1: alpha=0.89094, lambda=0.07584
## - Fold1: alpha=0.89094, lambda=0.07584
## + Fold1: alpha=0.03789, lambda=0.16332
## - Fold1: alpha=0.03789, lambda=0.16332
## + Fold1: alpha=0.26030, lambda=0.06467
## - Fold1: alpha=0.26030, lambda=0.06467
## + Fold2: alpha=0.26189, lambda=0.10504
## - Fold2: alpha=0.26189, lambda=0.10504
## + Fold2: alpha=0.17301, lambda=3.96064
## - Fold2: alpha=0.17301, lambda=3.96064
## + Fold2: alpha=0.89094, lambda=0.07584
## - Fold2: alpha=0.89094, lambda=0.07584
## + Fold2: alpha=0.03789, lambda=0.16332
## - Fold2: alpha=0.03789, lambda=0.16332

```

```

## + Fold2: alpha=0.26030, lambda=0.06467
## - Fold2: alpha=0.26030, lambda=0.06467
## + Fold3: alpha=0.26189, lambda=0.10504
## - Fold3: alpha=0.26189, lambda=0.10504
## + Fold3: alpha=0.17301, lambda=3.96064
## - Fold3: alpha=0.17301, lambda=3.96064
## + Fold3: alpha=0.89094, lambda=0.07584
## - Fold3: alpha=0.89094, lambda=0.07584
## + Fold3: alpha=0.03789, lambda=0.16332
## - Fold3: alpha=0.03789, lambda=0.16332
## + Fold3: alpha=0.26030, lambda=0.06467
## - Fold3: alpha=0.26030, lambda=0.06467
## + Fold4: alpha=0.26189, lambda=0.10504
## - Fold4: alpha=0.26189, lambda=0.10504
## + Fold4: alpha=0.17301, lambda=3.96064
## - Fold4: alpha=0.17301, lambda=3.96064
## + Fold4: alpha=0.89094, lambda=0.07584
## - Fold4: alpha=0.89094, lambda=0.07584
## + Fold4: alpha=0.03789, lambda=0.16332
## - Fold4: alpha=0.03789, lambda=0.16332
## + Fold4: alpha=0.26030, lambda=0.06467
## - Fold4: alpha=0.26030, lambda=0.06467
## + Fold5: alpha=0.26189, lambda=0.10504
## - Fold5: alpha=0.26189, lambda=0.10504
## + Fold5: alpha=0.17301, lambda=3.96064
## - Fold5: alpha=0.17301, lambda=3.96064
## + Fold5: alpha=0.89094, lambda=0.07584
## - Fold5: alpha=0.89094, lambda=0.07584
## + Fold5: alpha=0.03789, lambda=0.16332
## - Fold5: alpha=0.03789, lambda=0.16332
## + Fold5: alpha=0.26030, lambda=0.06467
## - Fold5: alpha=0.26030, lambda=0.06467
## Aggregating results
## Selecting tuning parameters
## Fitting alpha = 0.0379, lambda = 0.163 on full training set
print(df.glm)

## glmnet
##
## 42130 samples
##      55 predictor
##      2 classes: 'No', 'Yes'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 33704, 33705, 33703, 33704, 33704
## Resampling results across tuning parameters:
##
##     alpha      lambda      utility      prob_thr
##     0.03788688  0.16331812  1.0037390  0.3
##     0.17300505  3.96064125  0.1798124  0.1
##     0.26029747  0.06467323  0.6473321  0.3
##     0.26189424  0.10504447  0.3667833  0.3
##     0.89094150  0.07583706  0.1798124  0.1

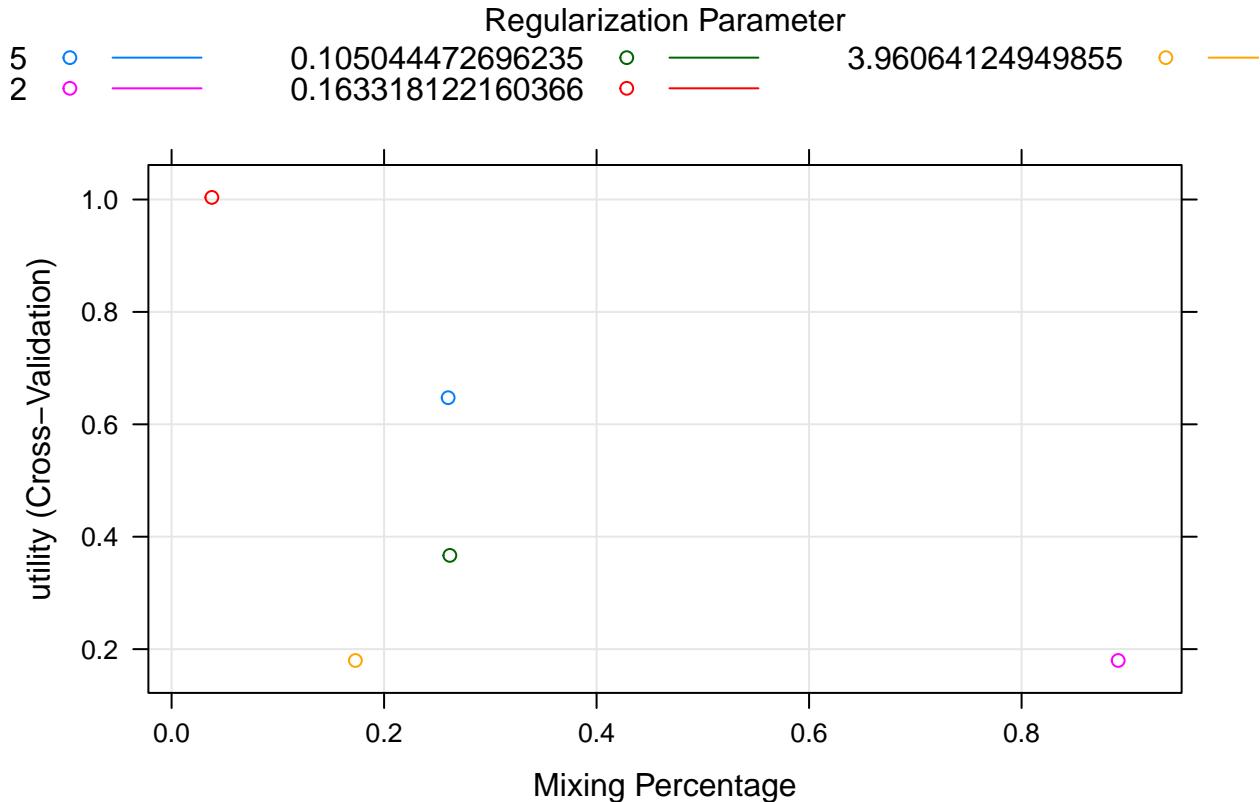
```

```

## utility was used to select the optimal model using the largest value.
## The final values used for the model were alpha = 0.03788688 and lambda
##   = 0.1633181.

plot(df.glm)

```



```

df.glm.model <- df.glm$finalModel

df.glm.model.coef <- predict(df.glm.model,
                               s = df.glm.model$lambdaOpt,
                               type = 'coefficients')

print(df.glm.model.coef)

## 84 x 1 sparse Matrix of class "dgCMatrix"
##                                     1
## (Intercept)           -8.527755e-01
## MonthlyRevenue        .
## MonthlyMinutes         -3.268115e-05
## TotalRecurringCharge -1.335488e-03
## DirectorAssistedCalls .
## OverageMinutes         2.390131e-04
## RoamingCalls          3.325798e-04
## PercChangeMinutes     -1.120819e-04
## PercChangeRevenues    2.058578e-04
## DroppedCalls          .
## BlockedCalls          .
## UnansweredCalls        .

```

```

## CustomerCareCalls      -2.203946e-03
## ThreewayCalls         -2.459976e-03
## ReceivedCalls          .
## OutboundCalls          .
## InboundCalls          -2.648910e-04
## PeakCallsInOut        -5.336858e-05
## OffPeakCallsInOut     -3.054776e-05
## DroppedBlockedCalls   .
## CallForwardingCalls   .
## CallWaitingCalls       .
## MonthsInService        .
## UniqueSubs             1.559815e-02
## ActiveSubs             .
## Handsets               .
## HandsetModels          -1.528450e-02
## CurrentEquipmentDays  3.776994e-04
## AgeHH1                 -1.287522e-03
## AgeHH2                 -1.138788e-04
## ChildrenInHHYes        8.992750e-03
## HandsetRefurbishedYes  1.011746e-01
## HandsetWebCapableYes   -1.106534e-01
## TruckOwnerYes          .
## RVOwnerYes              .
## HomeownershipUnknown    .
## BuysViaMailOrderYes    -1.241509e-02
## RespondsToMailOffersYes -2.458097e-02
## OptOutMailingsYes      .
## NonUSTravelYes          .
## OwnsComputerYes         .
## HasCreditCardYes        .
## RetentionCalls          1.660845e-01
## RetentionOffersAccepted .
## NewCellphoneUserYes     .
## NotNewCellphoneUserYes  .
## ReferralsMadeBySubscriber .
## IncomeGroup              .
## OwnsMotorcycleYes       .
## AdjustmentsToCreditRating -8.965122e-03
## HandsetPrice100          .
## HandsetPrice130          .
## HandsetPrice150          .
## HandsetPrice180          .
## HandsetPrice200          .
## HandsetPrice240          .
## HandsetPrice250          .
## HandsetPrice30            -5.287034e-04
## HandsetPrice300          .
## HandsetPrice40          .
## HandsetPrice400          .
## HandsetPrice500          .
## HandsetPrice60          .
## HandsetPrice80          .
## HandsetPriceUnknown      .
## MadeCallToRetentionTeamYes 2.168516e-01

```

```

## CreditRating2-High      2.650604e-02
## CreditRating3-Good     1.303344e-02
## CreditRating4-Medium   -1.928061e-02
## CreditRating5-Low      -1.388897e-01
## CreditRating6-VeryLow  .
## CreditRating7-Lowest    .
## PrizmCodeRural          1.850308e-03
## PrizmCodeSuburban        .
## PrizmCodeTown            .
## OccupationCrafts         .
## OccupationHomemaker      .
## OccupationOther           .
## OccupationProfessional    .
## OccupationRetired         .
## OccupationSelf             .
## OccupationStudent          .
## MaritalStatusUnknown      2.476788e-02
## MaritalStatusYes          .

df.glm.results <- fn_results(df.glm)

print('Umbral')

## [1] "Umbral"

print(df.glm.results$prob_thr)

## [1] 0.3

print('Utilidad en train')

## [1] "Utilidad en train"

print(df.glm.results$utility)

## [1] 1.003739

print('Utilidad en dev')

## [1] "Utilidad en dev"

df.glm.dev.prob <- predict(df.glm, newdata = df.part$dev, type = 'prob')
df.glm.dev.pred <- fn_pred(df.glm.dev.prob, thr = df.glm.results$prob_thr)

df.glm.dev.utility <- fn_utility(df.glm.dev.pred, df.part$dev$Churn)

print(df.glm.dev.utility)

## [1] 0.8547319

print('Matriz de confusion en dev')

## [1] "Matriz de confusion en dev"

df.glm.dev.cm <- conf_matrix(df.glm.dev.pred, df.part$dev$Churn)

print(df.glm.dev.cm)

##             Reference
## Prediction   No   Yes

```

```

##          No  2456  843
##          Yes   880  502
print('** Generacion de la prediccion sobre test sample')

## [1] "** Generacion de la prediccion sobre test sample"
test_sample <- read.csv('../data/test_sample.csv')
rownames(test_sample) <- test_sample$CustomerID
test_sample$CustomerID <- NULL
test_sample$ServiceArea <- NULL

file_id <- paste0(c(script.name, script.date), collapse = ' ')

gen_prediction(df.glm, test_sample, prob_thr = df.glm.results$prob_thr, id = file_id)

print('Done')

## [1] "Done"
script.done <- Sys.time()

print(script.done - script.start)

## Time difference of 1.779007 mins
# [1] "Utilidad real de la prediccion"
# [1] 0.9698145

```